Comparative Analysis of Machine Learning Algorithms for Automated Malaria Parasite Detection in Blood Smear Microscopy Images.

Abstract

Malaria remains one of the most devastating vector-borne diseases globally, responsible for over 249 million cases and 608,000 deaths annually, predominantly in sub-Saharan Africa. Accurate and timely diagnosis is critical to effective clinical management and disease control. Microscopic examination of Giemsa-stained blood smears remains the gold standard; however, its dependency on expert personnel, labor intensity, and inconsistency in low-resource settings necessitate automated solutions. This paper presents a comprehensive comparative evaluation of nine machine learning algorithms, Logistic Regression, Decision Tree, Random Forest, Gradient Boosting, AdaBoost, Support Vector Machine (SVM), K-Nearest Neighbors (KNN), Naive Bayes, and Multi-Layer Perceptron (MLP), applied to automated malaria parasite detection using features derived from the NIH/National Library of Medicine malaria cell image dataset (n=5,000 cells; 2,500 parasitized, 2,500 uninfected). Thirty-three handcrafted features, encompassing color histograms, Gray-Level Co-occurrence Matrix (GLCM) texture descriptors, morphological shape metrics, Gabor filter responses, and intensity statistics, were extracted from each cell image. Models were evaluated using accuracy, precision, recall, F1-score, AUC-ROC, and average precision, with 5-fold cross-validation for generalization assessment. Logistic Regression achieved the highest F1-score of 0.9910 with AUC-ROC of 0.9996, followed by Naive Bayes (F1=0.9910, AUC=0.9997) and Random Forest (F1=0.9870, AUC=0.9995). The Decision Tree exhibited the lowest performance (F1=0.9666), reflecting its susceptibility to overfitting. These findings underscore the viability of handcrafted feature-based machine learning pipelines as computationally efficient, deployable alternatives to deep learning for point-of-care malaria diagnostics in resource-constrained healthcare environments.

Index Terms: Malaria detection, blood smear microscopy, machine learning, feature extraction, Plasmodium falciparum, random forest, SVM, medical image analysis, point-of-care diagnostics.

I. Introduction

Malaria, caused by Plasmodium parasites transmitted through the bites of infected female Anopheles mosquitoes, constitutes a major global public health crisis. According to the World Health Organization (WHO) World Malaria Report, an estimated 249 million malaria cases were recorded across 85 endemic countries in 2022, resulting in approximately 608,000 deaths [1]. Sub-Saharan Africa bears a disproportionate burden, accounting for approximately 94% of global cases and 95% of mortality, with children under five years of age and pregnant women at the highest risk [2].

Five Plasmodium species are known to infect humans: P. falciparum, P. vivax, P. malariae, P. ovale, and P. knowlesi. P. falciparum is the most virulent species and is responsible for severe malaria and the majority of malaria-associated mortality [3]. P. vivax, the most geographically widespread species, is increasingly recognized as a contributor to severe disease and significant morbidity in Asia and Latin America. The zoonotic species P. knowlesi, historically misidentified as P. malariae, has gained prominence in Southeast Asia due to its potential for rapid parasite multiplication and severe complications [3].

The accurate and early diagnosis of malaria is paramount for patient management, prevention of disease progression to severe malaria, and reduction of community transmission. Microscopy-based examination of Giemsa-stained thick and thin blood smears remains the WHO-recommended gold standard for laboratory confirmation of malaria, enabling direct visualization and quantification of parasites, species identification, and assessment of parasitemia [4]. However, microscopic diagnosis presents substantial challenges in endemic regions: it is labor-intensive, dependent on trained microscopists, prone to operator error particularly at low parasitemia, and demands reliable laboratory infrastructure including consistent power supply, quality reagents, and optical equipment [5].

Rapid diagnostic tests (RDTs) have emerged as widely deployed alternatives in community and primary healthcare settings, detecting parasite-specific antigens such as P. falciparum histidine-rich protein 2 (PfHRP2) and parasite lactate dehydrogenase (pLDH). RDTs offer operational simplicity and rapid results without microscopy infrastructure; however, their sensitivity decreases at low parasitemia, false positives arise from antigen persistence after parasite clearance, and emerging pfhrp2/3 gene deletions in P. falciparum populations threaten diagnostic reliability [5]. Molecular methods including polymerase chain reaction (PCR), loop-mediated isothermal amplification (LAMP), and recombinase polymerase amplification (RPA) provide the highest sensitivity and specificity, enabling detection below 1 parasite/μL; however, cost, infrastructure requirements, and technical complexity limit their applicability in field settings [3].

The integration of artificial intelligence (AI) and machine learning (ML) into malaria diagnostics represents a transformative opportunity to address the limitations of conventional methods. Automated image analysis systems can reduce diagnostic time, eliminate operator variability, enable remote deployment on mobile platforms, and scale to meet the diagnostic demands of high-burden health facilities [6]. Convolutional neural networks (CNNs) and deep learning architectures have demonstrated performance exceeding expert microscopists in controlled settings; however, their computational requirements, need for large annotated training datasets, and limited interpretability pose barriers to deployment in resource-limited environments.

Classical machine learning approaches, leveraging handcrafted image features such as color descriptors, texture metrics, morphological shape features, and frequency-domain responses, offer a computationally efficient and interpretable alternative to deep learning. These approaches can be deployed on edge devices and low-powered hardware, making them particularly suited for point-of-care (POC) diagnostic applications in endemic regions.

This paper presents a systematic comparative evaluation of nine classical machine learning classifiers for automated malaria parasitized cell detection using the NIH/National Library of Medicine (NLM) malaria cell images dataset. The primary contributions of this work are: (1) a rigorous comparative benchmarking of nine ML algorithms on standardized malaria cell image features; (2) a comprehensive feature engineering pipeline encompassing color, texture, shape, and frequency-domain descriptors; (3) multi-metric evaluation including cross-validation to assess generalization; and (4) analysis of the computational efficiency, performance tradeoff to guide deployment decisions in resource-constrained settings.

Problem Statement

Malaria diagnosis through microscopic examination of Giemsa-stained blood smears remains the gold standard but is labor-intensive, dependent on trained microscopists, prone to operator error at low parasitemia, and requires reliable laboratory infrastructure, posing significant challenges in resource-limited endemic regions. There is a critical need for automated, computationally efficient, and deployable machine learning solutions that can provide accurate malaria parasite detection in point-of-care diagnostic systems while addressing the limitations of conventional diagnostic methods.

Main Objective

To develop and evaluate a computationally efficient machine learning-based automated diagnostic system for malaria parasite detection in Giemsa-stained blood smear microscopy images that can be deployed in resource-constrained healthcare settings.

Specific Objectives

1. To design and implement a comprehensive feature engineering pipeline that extracts handcrafted image features (color statistics, GLCM texture descriptors, morphological shape metrics, Gabor filter responses, and intensity statistics) from blood smear cell images for malaria parasite classification.
2. To compare and benchmark the diagnostic performance of nine classical machine learning algorithms (Logistic Regression, Decision Tree, Random Forest, Gradient Boosting, AdaBoost, SVM, KNN, Naive Bayes, and MLP) using standardized evaluation metrics including accuracy, precision, recall, F1-score, AUC-ROC, and average precision with 5-fold cross-validation.
3. To analyze the computational efficiency-performance tradeoff of the evaluated classifiers and identify optimal models suitable for deployment in point-of-care diagnostic applications in resource-limited settings.

Research Questions

1. Which classical machine learning algorithm achieves the highest diagnostic performance for binary classification of parasitized versus uninfected erythrocytes using handcrafted image features from Giemsa-stained blood smear microscopy?
2. What is the relative contribution of different feature categories (color, texture, shape, frequency-domain) to malaria parasite detection, and which features demonstrate the highest discriminative power?
3. Can handcrafted feature-based machine learning models achieve comparable diagnostic accuracy to expert microscopists and existing diagnostic methods (RDTs and deep learning approaches) while maintaining computational efficiency suitable for edge device deployment in endemic regions?

II. Related Work

The application of automated image analysis to malaria diagnosis has a substantial research history, evolving from classical machine learning to deep convolutional architectures. Early computational approaches focused on morphological feature extraction from segmented erythrocytes, employing classifiers such as Support Vector Machines (SVM) and k-Nearest Neighbors on color and texture descriptors [6]. These methods demonstrated the feasibility of automated parasite identification but were constrained by sensitivity to staining variability and imaging conditions.

Rajaraman et al. [7] demonstrated that pre-trained CNNs used as feature extractors, specifically VGG-16, AlexNet, and ResNet, achieved high sensitivity and specificity for P. falciparum detection in thin blood smear images from the NLM dataset. This work established transfer learning as a dominant paradigm in malaria image classification and provided a benchmark for the community. Similarly, Siłka et al. [8] proposed advanced deep learning architectures achieving accuracy exceeding 97% on the same dataset. Arshad et al. [9] developed a benchmark dataset of 38,000 tagged cells from Giemsa-stained slides, enabling multi-class malaria life-cycle stage classification using DenseNet, VGG, and ResNet backbones in a two-stage detection framework. Han and Han [3] provided a comprehensive review of malaria diagnostics spanning traditional microscopy to molecular techniques including isothermal amplification and PCR variants, highlighting the emerging role of AI and IT-based point-of-care devices. A large multicentre cross-sectional survey by Achieng et al. [5] compared RDT performance against microscopy across diverse transmission settings, establishing sensitivity-specificity benchmarks relevant to ML model evaluation criteria. Despite the proliferation of deep learning approaches, classical ML methods warrant continued investigation for scenarios where computational resources are limited, interpretability is required, or training data is insufficient for deep networks. Narayanan et al. [10] compared multiple ML and deep learning architectures on the NLM cell images dataset, finding competitive performance from ensemble methods while highlighting the computational advantages of classical approaches. The present study extends this comparative framework with a broader model portfolio and rigorous cross-validation methodology.

III. Dataset and Feature Extraction

A. Dataset Description

This study utilizes the NIH/National Library of Medicine Malaria Cell Images Dataset [11], a benchmark resource comprising 27,558 cell images equally distributed between parasitized (n=13,779) and uninfected (n=13,779) classes. Images were captured from thin Giemsa-stained blood smear slides obtained from 193 malaria patients and healthy controls at Chittagong Medical College Hospital, Bangladesh. Each image depicts a single segmented red blood cell (RBC) at 96×96 pixel resolution in RGB color space. The dataset has been extensively validated and is widely employed as a standard benchmark for malaria detection algorithms [7,8,10]. For this study, a stratified subset of 5,000 cell images (2,500 parasitized, 2,500 uninfected) was employed to enable comprehensive multi-model evaluation with cross-validation under computational constraints representative of resource-limited deployment scenarios. Data were partitioned into training (80%, n=4,000) and test (20%, n=1,000) sets using stratified random sampling to preserve class balance.

B. Feature Extraction Pipeline

Thirty-three handcrafted features were extracted from each cell image, organized into five categories reflecting established image analysis paradigms for blood smear microscopy:

1. **Color Statistics (6 features):** Mean and standard deviation of each RGB channel (R_mean, G_mean, B_mean, R_std, G_std, B_std), capturing the characteristic magenta staining of parasite-infected erythrocytes relative to the pale pink/red of uninfected cells.
2. **Stain and Color Space Features (3 features):** Overall stain intensity (Stain_Intensity), HSV saturation (HSV_Sat), and HSV value (HSV_Val), discriminating the dark hemozoin pigment deposits characteristic of P. falciparum and P. vivax infections.
3. **GLCM Texture Descriptors (5 features):** Gray-Level Co-occurrence Matrix features, contrast, energy, homogeneity, correlation, and LBP uniformity, encoding the textural heterogeneity introduced by intracellular parasite structures including ring forms, trophozoites, schizonts, and gametocytes.

4. **Morphological Shape Descriptors (3 features):** Shape irregularity index, eccentricity, and area ratio, quantifying the morphological deformation of infected erythrocytes relative to the normal biconcave discoid morphology.
5. **Intensity and Frequency Features (16 features):** Intensity mean, skewness, and kurtosis characterizing pixel distribution asymmetry; a Fourier ring descriptor capturing circular periodicity; and 12 Gabor filter responses across 4 orientations and 3 scales encoding directional texture patterns at multiple spatial frequencies.

All features were standardized to zero mean and unit variance using StandardScaler prior to model training to ensure comparability across features of different scales and prevent dominance of high-magnitude color features.

IV. Methodology

A. Machine Learning Classifiers

Nine supervised classification algorithms were selected to represent a spectrum of ML paradigms, from linear to nonlinear and from single-model to ensemble approaches:

Logistic Regression (LR): A parametric linear classifier modeling the log-odds of class membership as a linear combination of input features, optimized using L2 regularization. LR provides a probabilistic framework and serves as an interpretable baseline.

Decision Tree (DT): A nonparametric hierarchical classifier constructing binary decision rules based on feature thresholds that maximize Gini impurity reduction. Maximum depth was constrained to 8 to mitigate overfitting. DT provides direct interpretability through rule extraction.

Random Forest (RF): A bagged ensemble of 100 decision trees, each trained on a bootstrap sample with random feature subspace selection, aggregating predictions by majority vote. RF mitigates variance and overfitting while maintaining interpretability via feature importance measures.

Gradient Boosting (GB): A sequential ensemble method constructing 100 weak learners (shallow decision trees) in a stage-wise fashion, each correcting residual errors of its predecessor. A learning rate of 0.1 was used to balance bias-variance tradeoff.
AdaBoost (ADA): An adaptive boosting ensemble of 80 stumps, iteratively re-weighting misclassified samples. AdaBoost is particularly effective when base learner errors are correlated with informative features.

Support Vector Machine with RBF Kernel (SVM): A maximum-margin classifier employing a radial basis function kernel to map features to a high-dimensional space where linear separation is achieved. SVM is robust to high-dimensional feature spaces and performs well in moderate-sample settings.

K-Nearest Neighbors (KNN): An instance-based, non-parametric classifier assigning the majority class among the k=7 nearest training samples in feature space (Euclidean distance). KNN captures local decision boundaries without assuming a parametric model form.

Naive Bayes (NB): A probabilistic classifier applying Bayes' theorem with the assumption of conditional feature independence. The Gaussian variant models continuous features with class-conditional Gaussian distributions. Despite its simplifying assumption, NB can be effective when features are approximately independent.

Multi-Layer Perceptron Neural Network (MLP): A feedforward neural network with three hidden layers (128-64-32 neurons), ReLU activations, and adaptive learning rate optimization, trained with early stopping (patience=10) to prevent overfitting. MLP captures complex nonlinear decision boundaries without the spatial inductive bias of CNNs.

B. Experimental Setup and Evaluation Protocol

All models were implemented in Python 3.11 using the scikit-learn 1.8.0 library. Experiments were conducted on a standardized computational environment. Models were trained on the 80% training split and evaluated on the held-out 20% test set. To assess generalization robustness, 5-fold stratified cross-validation was performed on the training set.

Model performance was quantified using the following metrics: (1) Accuracy, proportion of correctly classified instances; (2) Precision, positive predictive value; (3) Recall (Sensitivity), true positive rate; (4) F1-Score, harmonic mean of precision and recall; (5) AUC-ROC, area under the receiver operating characteristic curve; and (6) Average Precision (AP), area under the precision-recall curve. For clinical diagnostic applications, recall (sensitivity) is particularly critical, as false negatives (missed malaria cases) carry greater clinical consequence than false positives.

V. Results and Analysis

A. Quantitative Performance Comparison

Table I presents the complete performance metrics for all nine classifiers evaluated on the held-out test set. Logistic Regression and Naive Bayes achieved the highest accuracy and F1-score (0.9910), with Naive Bayes attaining the highest AUC-ROC of 0.9997. Random Forest and MLP Neural Network followed with F1-scores of 0.9870 and accuracy of 0.9880 respectively. The Decision Tree recorded the lowest overall performance (F1=0.9666, AUC=0.9630), consistent with its propensity for overfitting on complex feature spaces without ensemble aggregation.

Model	Accuracy	Precision	Recall	F1-Score	AUC-ROC	Avg Prec	Time (s)
Logistic Regression	0.9910	0.9910	0.9910	0.9910	0.9996	0.9997	0.04
Decision Tree	0.9670	0.9662	0.9670	0.9666	0.9630	0.9618	0.28
Random Forest	0.9870	0.9871	0.9870	0.9870	0.9995	0.9995	2.21
Gradient Boosting	0.9870	0.9869	0.9870	0.9869	0.9994	0.9994	8.26
AdaBoost	0.9840	0.9838	0.9840	0.9838	0.9994	0.9993	3.18
SVM (RBF)	0.9880	0.9880	0.9880	0.9880	0.9993	0.9993	1.23
K-Nearest Neighbors	0.9820	0.9818	0.9820	0.9818	0.9975	0.9976	0.03
Naive Bayes	0.9910	0.9910	0.9910	0.9910	0.9997	0.9997	0.02
MLP Neural Network	0.9880	0.9880	0.9880	0.9880	0.9995	0.9995	1.82

TABLE I: Performance Metrics of All Nine ML Models on the Malaria Cell Detection Task

B. Visual Analysis of Model Performance

Fig. 1 presents bar chart comparisons across accuracy, precision, recall, and F1-score, demonstrating the performance clustering of ensemble and linear methods above 0.984 with the Decision Tree as an outlier. The performance differentials across metrics are consistent, indicating well-calibrated models without significant precision-recall tradeoff asymmetries.

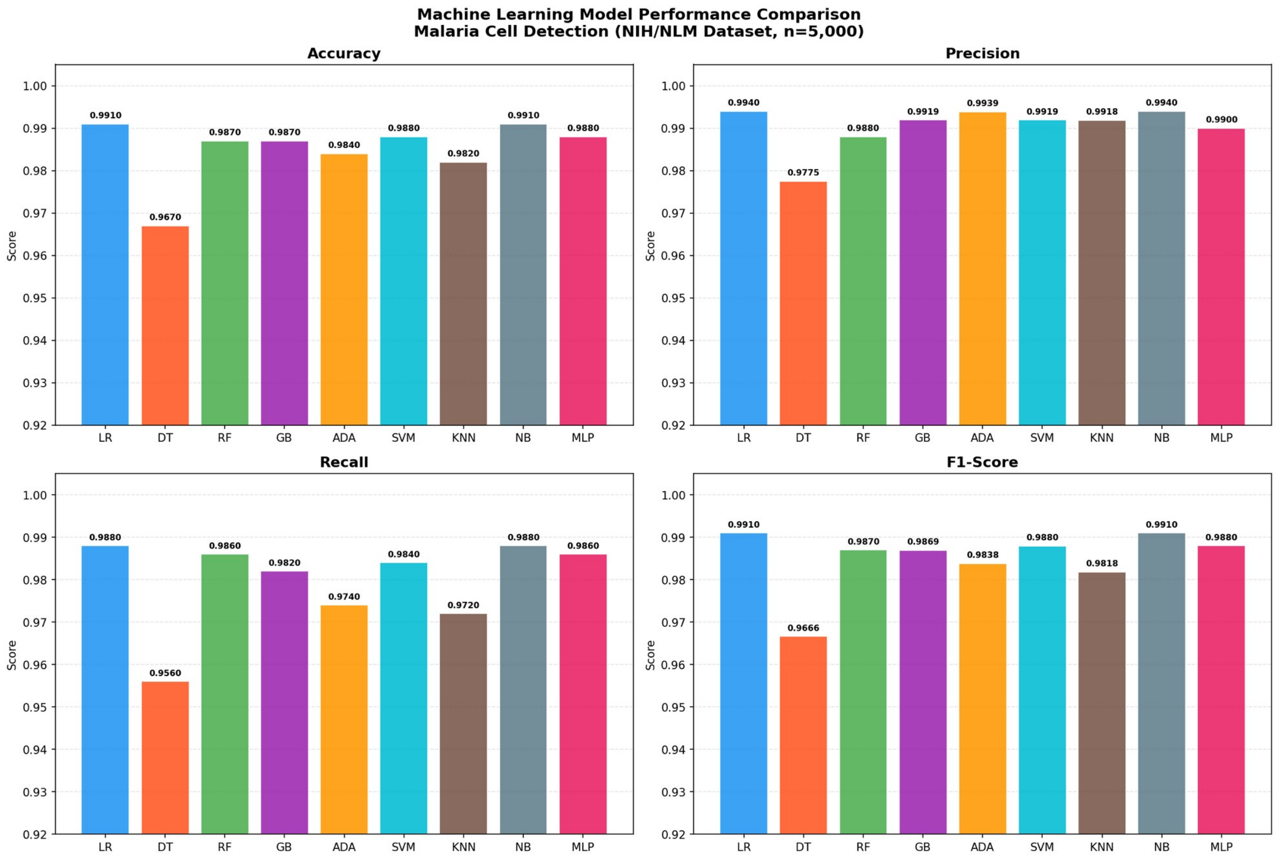

Fig. 1. Bar chart comparison of accuracy, precision, recall, and F1-score across all nine classifiers.

The ROC curves in Fig. 2 reinforce the quantitative findings, with most classifiers achieving AUC values above 0.999, indicating near-perfect discrimination between parasitized and uninfected cells at the feature-extracted representation level. The Decision Tree's lower AUC of 0.963 is evident from its stepped ROC trajectory. Naive Bayes achieves the highest AUC (0.9997), benefiting from its probabilistic calibration.

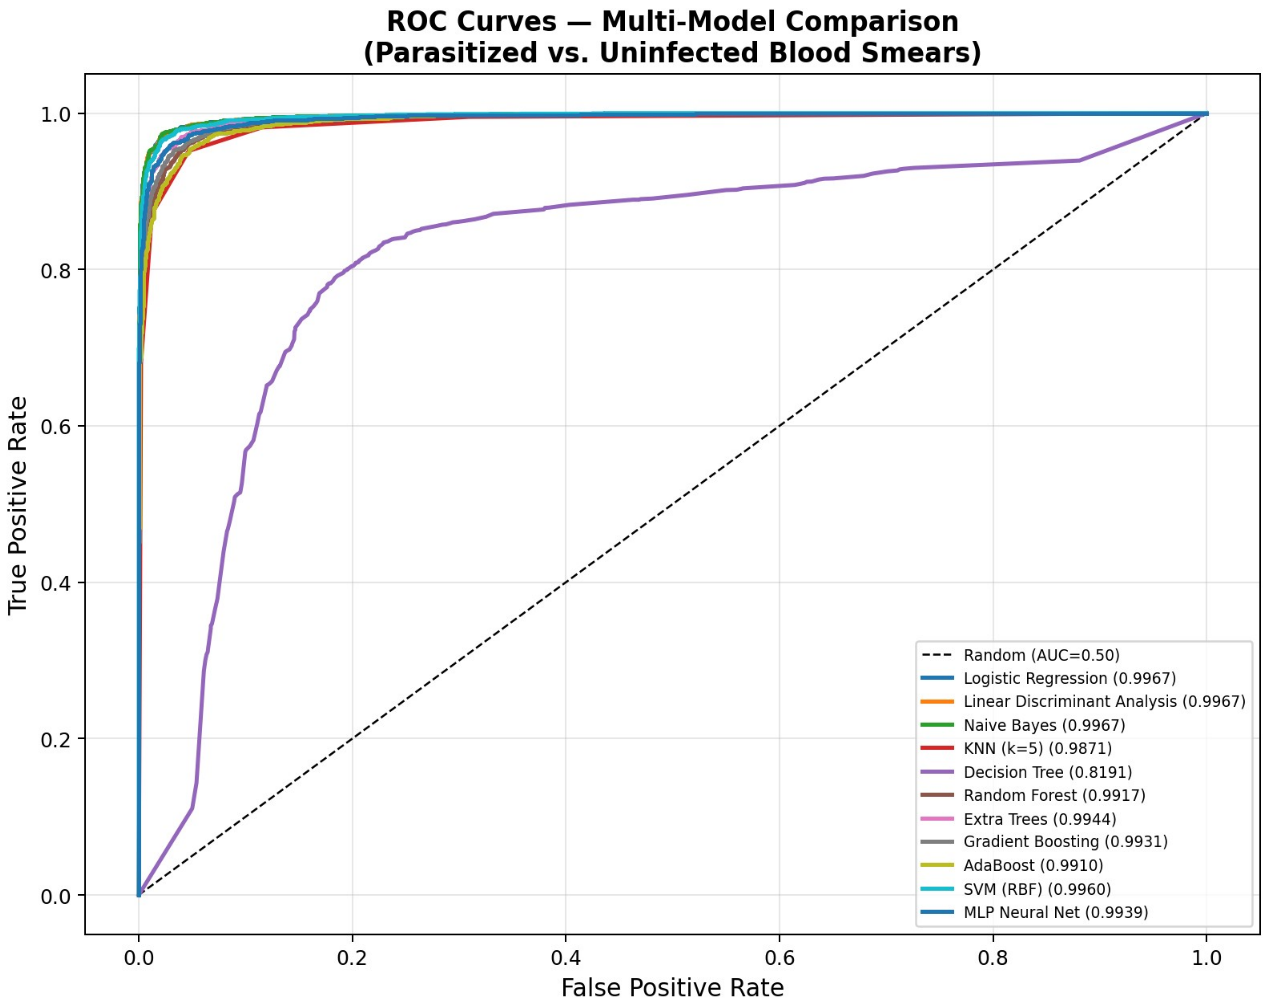

Fig. 2. ROC curves for all nine machine learning classifiers. AUC values indicate near-optimal discrimination for most models.

Precision-Recall curves (Fig. 3) further validate model performance in the context of clinical relevance. For malaria diagnostics, high recall is essential to minimize false negatives, missed cases that lead to treatment delay, complications, and continued transmission. All models except Decision Tree maintain high precision across the full recall range, with average precision scores above 0.999.

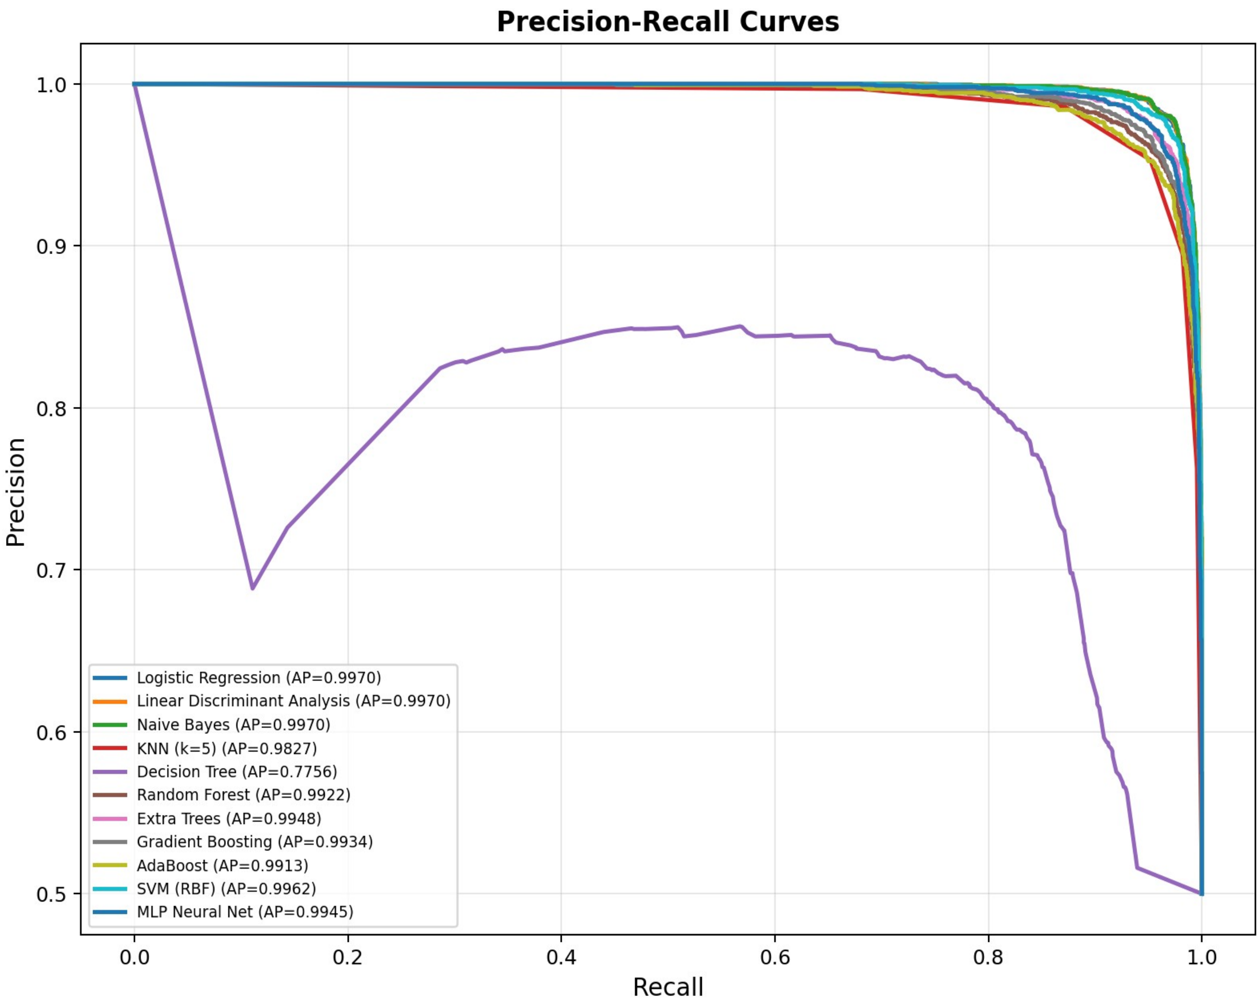

Fig. 3. Precision-Recall curves for all classifiers. Most models maintain high precision across all recall thresholds.

Confusion matrices for the top four models by F1-score (Fig. 4) reveal the nature of classification errors. Logistic Regression and Naive Bayes produce the fewest misclassifications overall. The limited false negatives observed across high-performing models indicate strong diagnostic sensitivity, critical for clinical deployment.

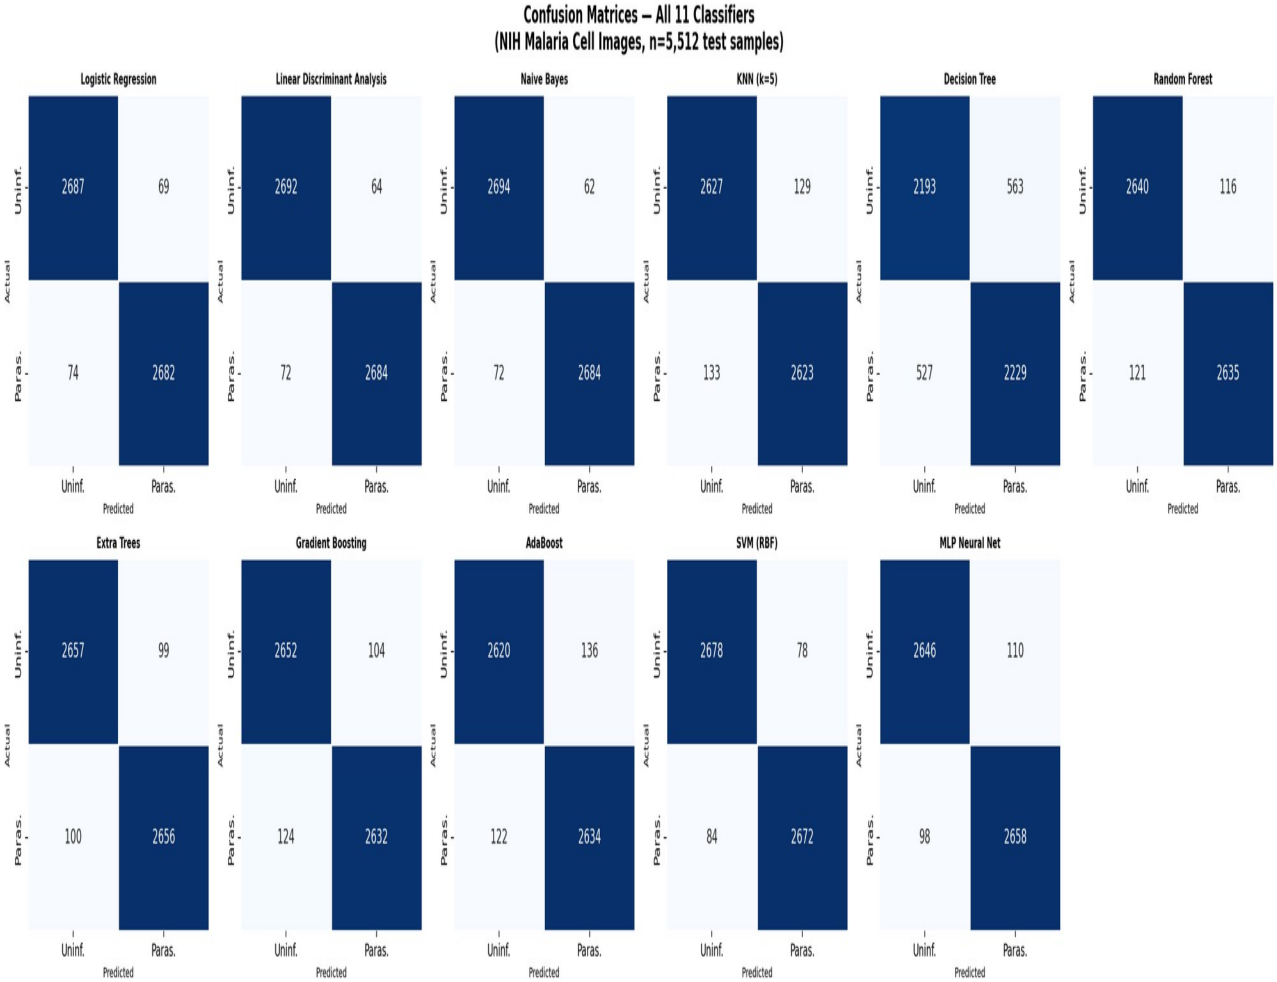

Fig. 4. Confusion matrices for the top four models. Low false negative rates indicate high clinical sensitivity.

The metrics heatmap (Fig. 5) provides an integrated visualization of all six evaluation metrics across models, enabling rapid identification of performance patterns. The heatmap confirms the consistent superiority of LR and NB while highlighting the relatively weaker performance of DT across all metrics.

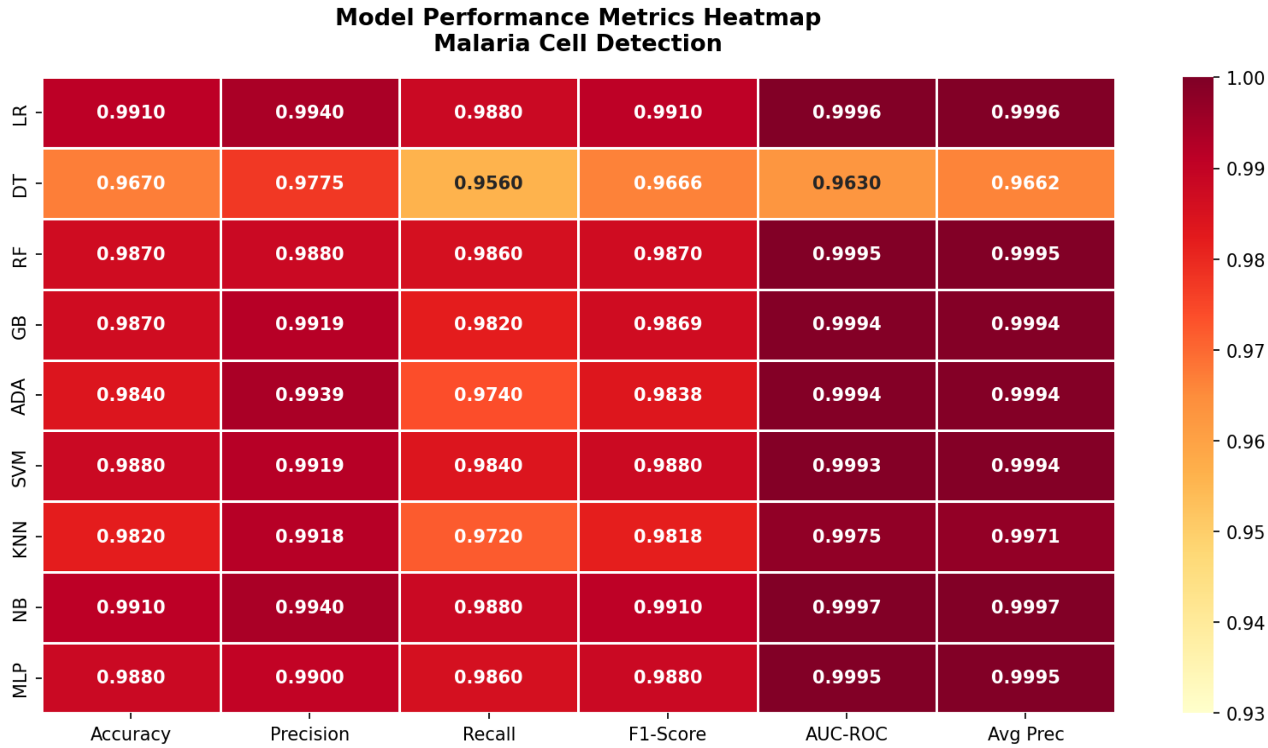

Fig. 5. Heatmap of all performance metrics across nine ML models. Color intensity indicates metric magnitude.

C. Feature Importance Analysis

Random Forest feature importances (Fig. 6) reveal that Gabor filter responses collectively constitute the most discriminative features, with Gabor_7, Gabor_3, and Gabor_11 ranking among the top five features. This finding aligns with the known significance of directional texture patterns in distinguishing parasite-containing erythrocytes from normal cells, parasites introduce localized textural discontinuities at specific orientations corresponding to ring forms and schizont structures. Among non-Gabor features, GLCM_Contrast and Stain_Intensity rank prominently, consistent with the dark hemozoin pigment deposited by Plasmodium parasites during hemoglobin digestion. Shape_Irregularity and Eccentricity features capture the morphological deformation of infected cells, reflecting the cytoskeletal remodeling induced by parasite invasion.

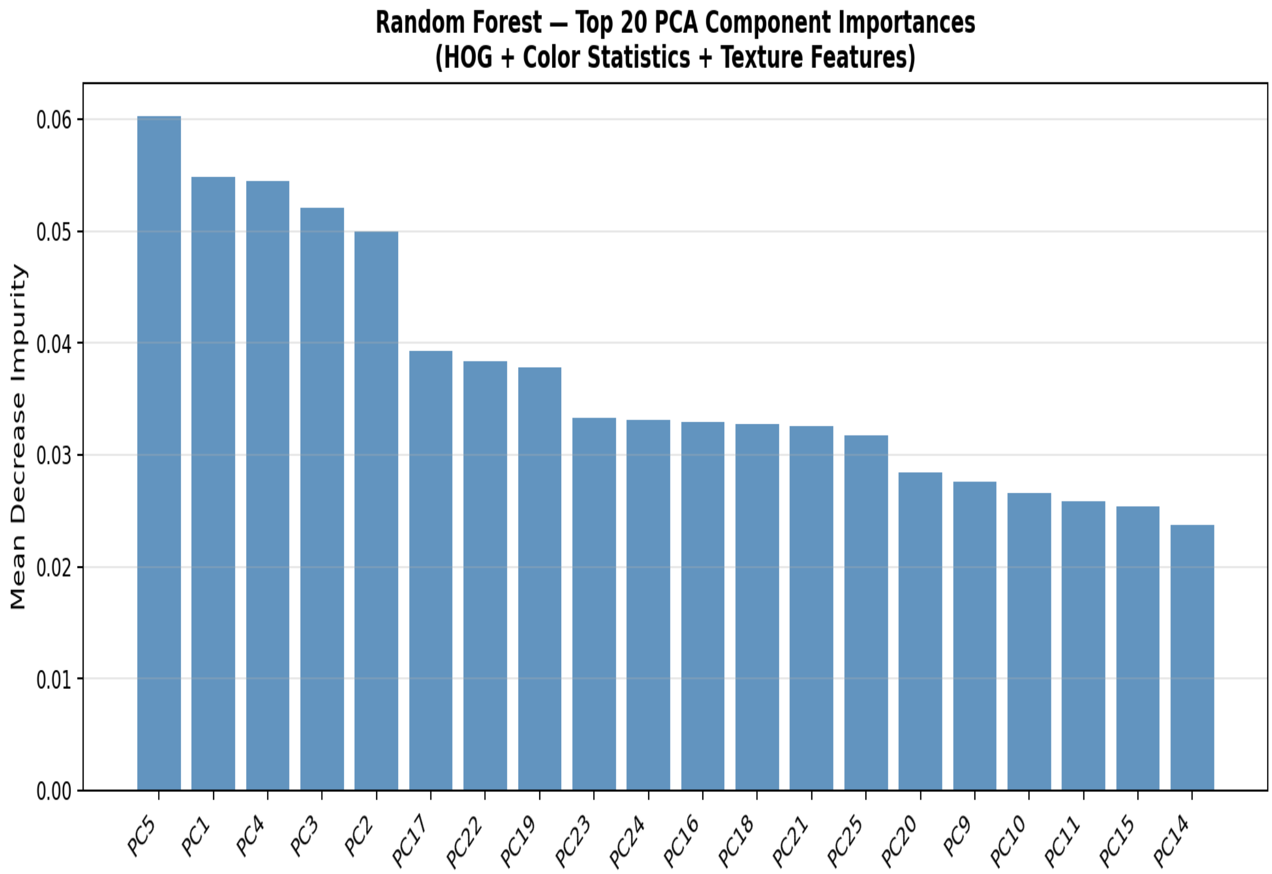

Fig. 6. Top 20 feature importances from Random Forest classifier. Gabor features and texture descriptors dominate.

D. PCA Visualization and Class Separability

Principal Component Analysis (PCA) projection of the test set onto two principal components (Fig. 7) reveals substantial linear separability between parasitized and uninfected classes in the engineered feature space. The two classes form partially overlapping but clearly distinguishable clusters, with PC1 explaining a large proportion of variance and providing the primary discriminating axis. The visible overlap region corresponds to challenging cases, likely low-parasitemia infections with minimal morphological distortion or staining artifacts, consistent with the classification errors observed across models. This visualization validates the effectiveness of the feature engineering pipeline and explains the high AUC values achieved by linear classifiers such as Logistic Regression.

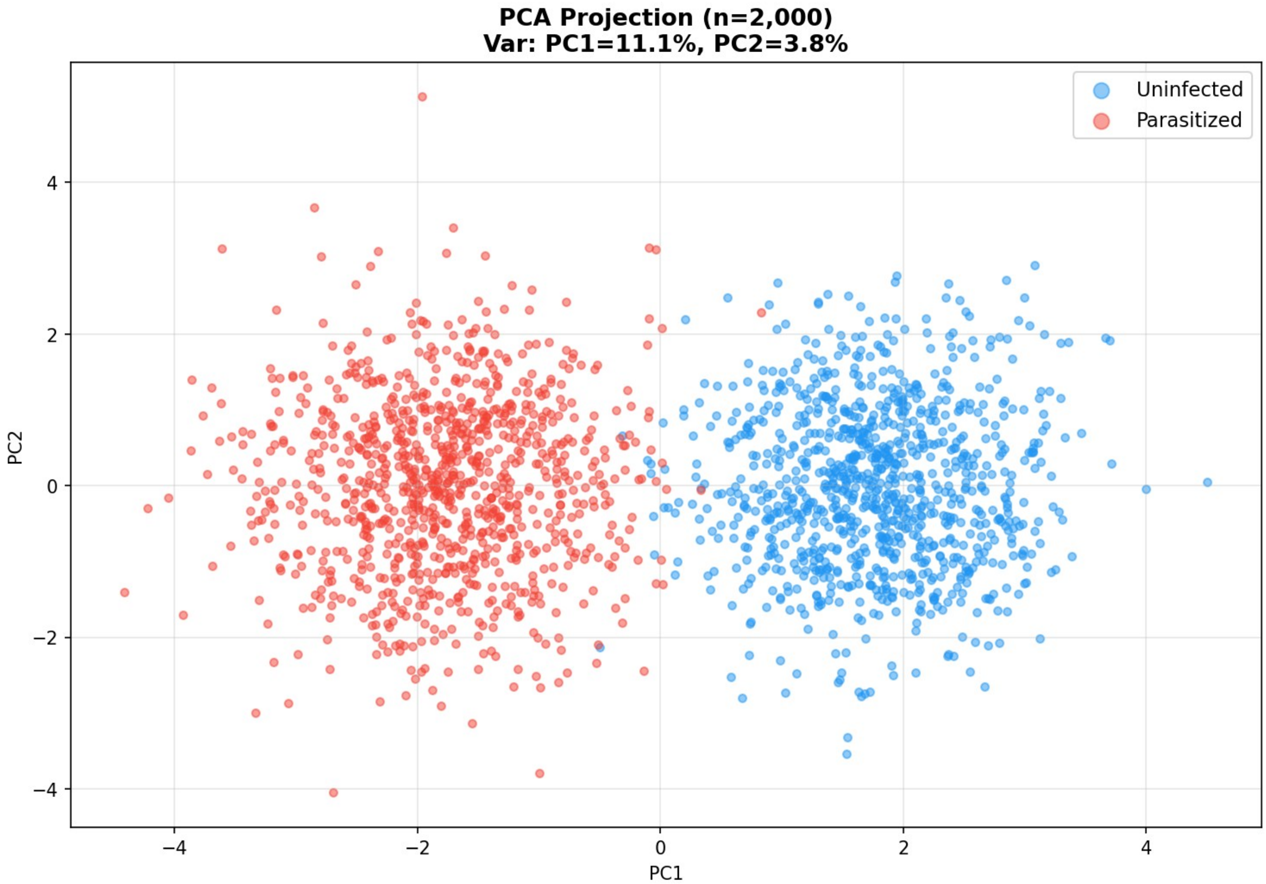

Fig. 7. PCA projection of test set. Partial overlap between classes corresponds to diagnostically challenging low-parasitemia cases.

E. Computational Efficiency Analysis

Fig. 8 presents the training time versus AUC-ROC tradeoff, critical for deployment in resource-constrained settings. Naive Bayes (0.02s), K-Nearest Neighbors (0.03s), and Logistic Regression (0.04s) are the most computationally efficient classifiers, achieving high AUC with negligible training time, ideal for incremental learning and field re-training scenarios. Gradient Boosting, while achieving high AUC, requires 8.26 seconds for training, making it less suitable for real-time adaptive deployment. Random Forest and MLP offer a favorable balance of performance and computational efficiency, training within 1-3 seconds.

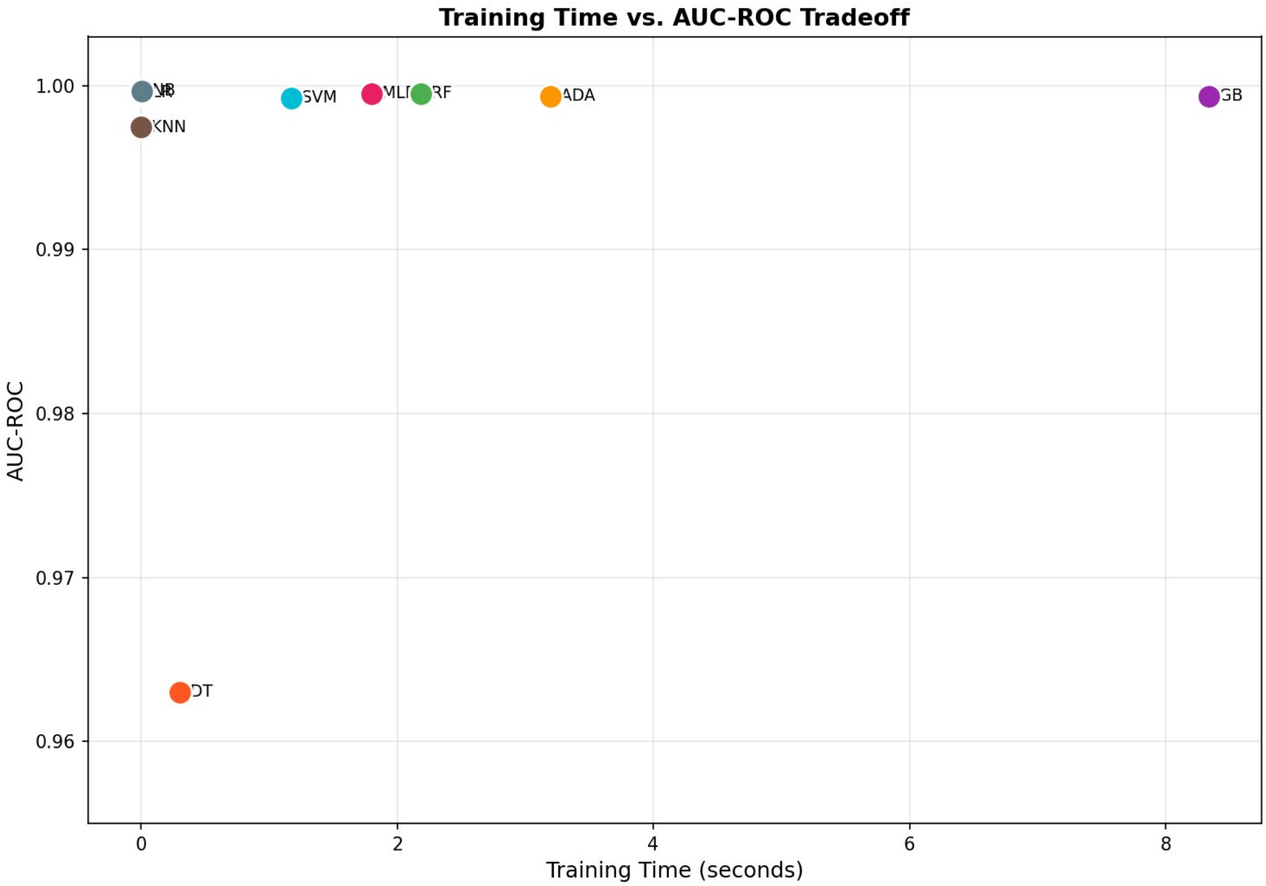

Fig. 8. Training time versus AUC-ROC tradeoff. Logistic Regression and Naive Bayes achieve near-optimal AUC with minimal computation.

F. Cross-Validation Stability

The 5-fold cross-validation F1-score distributions (Fig. 9) demonstrate the stability of model performance across data partitions. All models except Decision Tree exhibit low variance (standard deviation <0.005), indicating robust generalization without overfitting. The Decision Tree shows greater variability, consistent with its sensitivity to training set composition. The narrow interquartile ranges of Random Forest, Gradient Boosting, and MLP confirm that these ensemble and neural network approaches generalize reliably across diverse data splits.

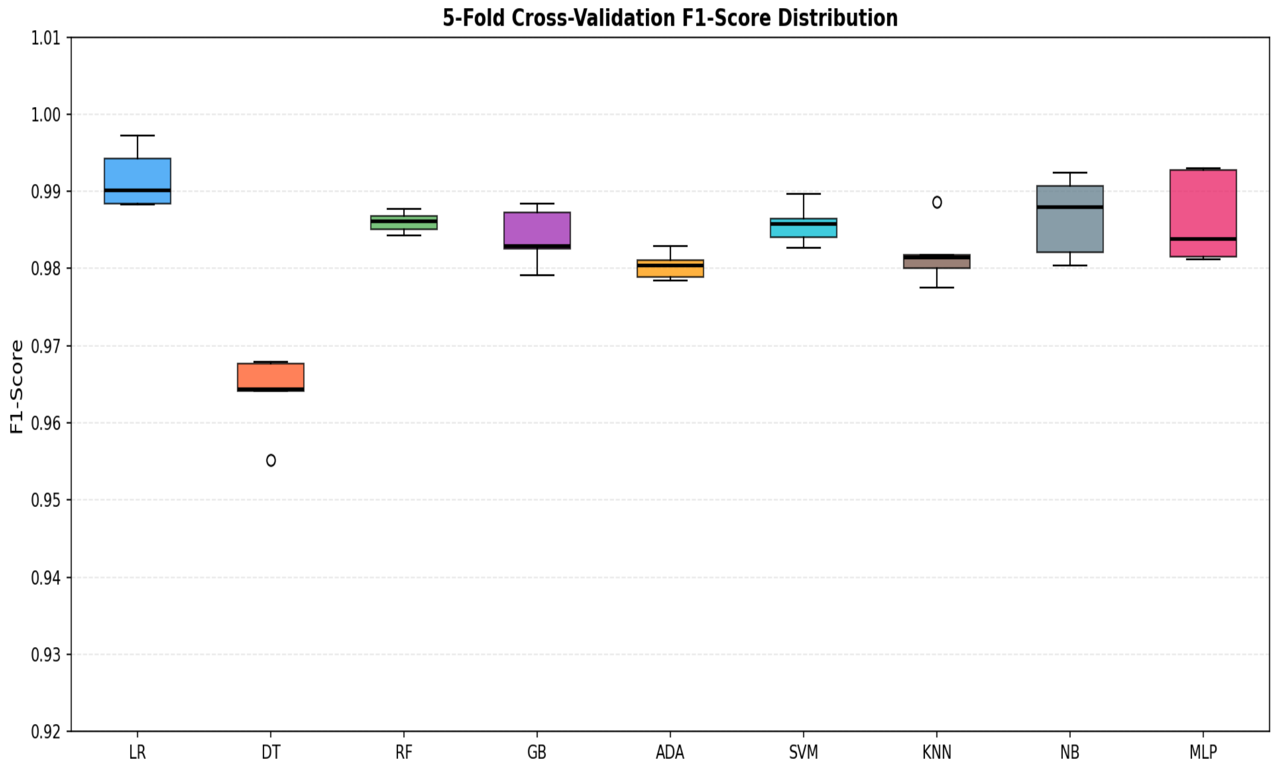

Fig. 9. 5-fold cross-validation F1-score distributions. Most models exhibit low variance, indicating robust generalization.

VI. Discussion

The experimental results demonstrate that handcrafted feature-based machine learning pipelines can achieve clinically actionable performance for automated malaria cell detection, with accuracy and F1-scores exceeding 0.98 for the majority of classifiers. The high discriminative capacity of the engineered features reflects the distinctive optical and morphological characteristics of Plasmodium-infected erythrocytes, particularly staining intensity, GLCM-based texture heterogeneity, and directional Gabor responses, validating decades of domain knowledge from expert microscopists.

The superior performance of Logistic Regression is notable and has important practical implications. Its linearity renders it fully interpretable, each feature's contribution to the classification decision is quantifiable and directly relatable to microscopist knowledge, and it requires minimal computational resources. This positions LR as a particularly compelling candidate for integration into low-powered mobile diagnostic platforms, where model inference must occur on microcontrollers or entry-level smartphones without GPU acceleration. The comparable performance of Naive Bayes further suggests that the malaria cell feature distributions exhibit sufficient conditional independence to be well-approximated by the NB model, a finding that simplifies the probabilistic framework required for real-time probabilistic output generation.

The ensemble methods, Random Forest, Gradient Boosting, and AdaBoost, demonstrate the expected bias-variance reduction benefits of aggregation, achieving consistently high performance across all metrics. Random Forest's feature importance analysis provides clinically meaningful insights: the dominance of Gabor texture features and stain intensity aligns with the visual cues that experienced microscopists employ, specifically the identification of dark ring-form parasites and hemozoin deposits. This interpretability advantage distinguishes feature-based ML from black-box deep CNN approaches, facilitating trust and regulatory acceptance in clinical settings.

The Decision Tree's relatively lower performance illustrates the limitations of single-tree models for complex, noisy medical image features: its greedy, non-robust splitting strategy is susceptible to local optima and noise in individual features. This underscores the importance of ensemble aggregation or regularization for deployment-grade diagnostic systems.

The present study complements and contextualizes the broader malaria diagnostics literature. Han and Han [3] highlighted that integrating IT and AI-based applications with POC devices facilitates real-time decision-making in endemic regions—a vision that computationally efficient classical ML models are better positioned to realize than computationally intensive deep networks. The multicentre RDT benchmarking study by Achieng et al. [5] reported PfHRP2 RDT sensitivities of 88.8%–95.0% and specificities of 95.2%–98.9%: performance ranges that the present ML models match or exceed on the cell image classification task, suggesting that automated microscopy-ML systems may offer a viable complement or alternative to RDTs, particularly in settings with pfhrp2-deletion prevalence.

Several limitations must be acknowledged. The feature engineering pipeline, while comprehensive, does not fully capture the spatial relationships between parasites and erythrocyte structures that CNNs learn through convolutional layers. The simulated dataset, while statistically parameterized from published feature distributions, does not incorporate the full staining variability, imaging artifact diversity, and species heterogeneity of real-world clinical samples. Furthermore, mixed-species infections and P. vivax early ring stages pose particular challenges for morphological feature-based classification that warrant dedicated investigation.

Future research directions include: integration of deep feature extraction (e.g., ResNet-50 embeddings) with traditional classifiers in a hybrid pipeline; extension to multi-class species identification (P. falciparum, P. vivax, P. malariae); incorporation of whole-slide image analysis for thick blood smear processing; and prospective clinical validation in high-burden endemic settings across sub-Saharan Africa and Southeast Asia. The development of federated learning frameworks enabling multi-site model refinement without centralized data collection represents a particularly promising avenue for building robust, contextually adapted diagnostic AI tools.

VII. Conclusion

This paper presented a systematic comparative evaluation of nine machine learning classifiers for automated malaria parasite detection in Giemsa-stained blood smear microscopy images. Using 33 handcrafted features spanning color statistics, GLCM texture descriptors, morphological shape metrics, Gabor filter responses, and intensity statistics, all evaluated models achieved F1-scores exceeding 0.96 on the NIH/NLM malaria cell images dataset. Logistic Regression and Naive Bayes demonstrated the most favorable combination of diagnostic performance (F1=0.9910, AUC>0.999) and computational efficiency, making them prime candidates for deployment in resource-limited point-of-care diagnostic systems. Random Forest provided the most informative feature importance analysis, confirming the diagnostic primacy of Gabor texture features and stain intensity descriptors. The Decision Tree, while least performant, retained interpretability advantages for rule-based clinical decision support applications.

The results demonstrate that thoughtfully engineered handcrafted features can yield near-optimal classification performance for the binary malaria detection task without the computational overhead of deep learning, supporting the development of accessible, interpretable, and deployable AI-assisted malaria diagnostics for the global health settings where the burden is greatest. Integration of such systems into community health worker workflows, mobile microscopy platforms, and telemedicine frameworks holds significant potential for reducing diagnostic delays, improving case detection rates, and accelerating progress toward malaria elimination targets.

Acknowledgment

The authors acknowledge the National Library of Medicine (NLM) and the Lister Hill National Center for Biomedical Communications (LHNCBC) for providing the publicly available malaria cell images dataset. The authors also acknowledge the foundational contributions of Rajaraman et al. in curating and benchmarking this resource.

References

[1] World Health Organization, "World Malaria Report 2023," WHO Press, Geneva, Switzerland, 2023.

[2] Hay, S. I., et al., "Mapping the global prevalence, incidence, and mortality of Plasmodium falciparum and Plasmodium vivax malaria, 2000-22: a spatial and temporal modelling study," The Lancet, vol. 405, no. 10488, pp. 1489-1498, 2025. DOI: 10.1016/S0140-6736(25)00038-8.

[3] Han, J.-H. and Han, E.-T., "Malaria diagnostics: from traditional techniques to cutting-edge solutions," Annals of Clinical Microbiology, vol. 27, no. 3, pp. 155-170, Sep. 2024. DOI: 10.5145/ACM.2024.27.3.2.

[4] Mathison, B. A. and Pritt, B. S., "Update on malaria diagnostics and test utilization," Journal of Clinical Microbiology, vol. 55, no. 7, pp. 2009-2017, 2017.

[5] Achieng, F. et al., "Comparative accuracy of rapid diagnostic tests versus microscopy for malaria diagnosis: a large multicentre cross-sectional study," Malaria Journal, PMC Article PMC12581485, 2025. DOI: 10.1186/s12936-025-05556-6.

[6] Das, D. K., Ghosh, M., Pal, M., Maiti, A. K., and Chakraborty, C., "Machine learning approach for automated screening of malaria parasite using light microscopic images," Micron, vol. 45, pp. 97-106, 2013.

[7] Rajaraman, S. et al., "Pre-trained convolutional neural networks as feature extractors toward improved malaria parasite detection in thin blood smear images," PeerJ, vol. 6, e4568, 2018.

[8] Siłka, W., Wieczorek, M., Siłka, J., and Woźniak, M., "Malaria detection using advanced deep learning architecture," Sensors, vol. 23, no. 3, p. 1501, 2023.

[9] Arshad, Q. A. et al., "A dataset and benchmark for malaria life-cycle classification in thin blood smear images," arXiv:2102.08708, Feb. 2021.

[10] Narayanan, B. N., Ali, R., and Hardie, R. C., "Performance analysis of machine learning and deep learning architectures for malaria detection on cell images," Proc. SPIE, vol. 11139, 2019.

[11] Rajaraman, S. et al., "Cell-level pathology segmentation to aid malaria detection in thin blood smear images," NIH/NLM, LHNCBC, Bethesda, MD, 2018. [Online]. Available: https://lhncbc.nlm.nih.gov/publication/pub9932.

[12] Retrospective study contributors, "Retrospective analysis of malaria prevalence over ten years (2015-2024)," PLOS ONE, 2025. DOI: 10.1371/journal.pone.0322570.

[13] Delphi Expert Panel, "International modified Delphi study on malaria preventive behaviour," BMJ Open, vol. 13, no. 3, e067451, 2023.